

*A Comparison of Sparse and EBE Storage
Schemes on the Efficiency of Parallel
Conjugate Gradient Iterative Methods for
Finite Element Analysis.*

Shivaraju Gowda

Date: July 18, 2002

Advisor: Dr. Lonny L Thompson

Department of Mechanical Engineering,
Clemson University



Computational Mechanics Laboratory



Summary

- **Preconditioned conjugate gradient iterative** method for FEA in parallel on distributed memory machines using MPI and Fortran90.
- Comparing two types of matrix storage schemes

EBE – Element-By-Element

CSR – Compressed Sparse Row

- Implemented **Jacobi Preconditioner**.
- Scalability analysis on two type of parallel machines.
- Applied to structured and unstructured mesh for 2D and 3D problems.



Outline

1. Introduction
 - Motivation
2. Formulation and Implementation
 - Finite Element Formulation
 - Matrix Storage Schemes
 - Conjugate Gradient Algorithm
 - Preconditioners
3. Parallel Implementation
 - Domain Decomposition
 - Parallel CG Algorithm
 - Implementation
4. Numerical Examples
 - Poisson's Equation – with Dirichlet BC
 - Scalability Analysis
 - Unstructured Mesh
5. Conclusions



Introduction


Why Parallel ?

Use of Multiple Processors to process single application simultaneously


- Reduce total wall clock simulation time
- Solve large complex problems which do not fit into the local memory of a single system

Architectures:

Shared Memory (OpenMP, HPF)



Distributed memory(MPI)



Introduction

Why Iterative Methods ?



- **Direct Methods**, robust and predictable. For 3D problems, Computational time and memory requirements increase with $O(n^3)$.
- **Iterative methods**, suitable for such problems, especially for sparse matrices from finite element discretization.
- Iterative methods use the knowledge of the problem for better convergence
- Suitable for parallel implementation.
- With good preconditioner can match the performance of Direct methods.



- No general purpose routine for all problems.
- Knowledge of problems is required on the part of the end user to select the best method, preconditioner and initial value.



Finite Element Formulation

Second order linear elliptic partial differential equation

$$-\underline{\nabla} \cdot (a(\underline{x}) \underline{\nabla} u) = f(\underline{x}), \quad \underline{x} \in \Omega \subset \Re^2$$

subject to Dirichlet Boundary Condition

$$u|_{\partial\Omega} = 0$$

The domain Ω is discretized into a non-overlapping set of elements.

Piecewise linear finite element approximation,

$$u = \sum_{i=1}^N u_i N_i(\underline{x}), \quad v = \sum_{i=1}^N v_i N_i(\underline{x})$$

where $N_i(\underline{x})$ are piecewise linear basis functions for $i = 1, \dots, N$

$$\int_{\Omega} a \underline{\nabla} v \cdot \underline{\nabla} u d\underline{x} = \int_{\Omega} f v d\underline{x}$$



Matrix Problem

The set of equations may be written in matrix notation as

$$\underline{A}\underline{u} = \underline{b}$$

Where,

$$\underline{u} = (u_1, \dots, u_N)^T$$

$$A_{ji} = \int_{\Omega} a \underline{\nabla} N_j \cdot \underline{\nabla} N_i d\underline{x},$$

$$b_j = \int_{\Omega} f N_j d\underline{x}$$


General solution Procedure :

- Calculate the stiffness coefficients of all the elements.
- Assemble the global stiffness matrix A .
- Solve the system of equations using an iterative algorithm.




Matrix Storage


Full storage




Skyline storage



Sparse storage



EBE storage



Skyline Storage

- Stores all coefficients within skyline, including zeros
- Requires renumbering to minimize skyline
- Efficient for direct solvers on shared-memory machines

Sparse Storage

- Stores nonzero coefficients only
- Predominantly indirect access patterns
- Efficient for iterative methods on distributed memory architectures

Element by Element storage

- Retains only element-level matrices, no global assembly
- Efficient for iterative methods on distributed memory
- Matrix-vector product done in 3 stages, gather to element local, matrix-vector product at element level and scatter to global
- Element matrices can be recalculated as needed minimal storage but slower



Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

Matrix Storage Schemes

Element by Element (EBE)

- Originally developed to reduce memory.
- Provides significant opportunities for fine-grained (vectorized) parallel computation.
- Simpler data structure.
- Requires the preconditioners to be expressed in element level calculation, limiting choices.

$$AA = [1. \ 2. \ 3. \ 4. \ 5. \ 6. \ 7. \ 8. \ 9. \ 10. \ 11. \ 12.]$$

$$JA = [1 \ 4 \ 1 \ 2 \ 4 \ 1 \ 3 \ 4 \ 5 \ 3 \ 4 \ 5]$$

$$IA = [1 \ 3 \ 6 \ 10 \ 12 \ 13]$$

Compressed Sparse Row (CSR)

- Lower Memory Requirements $O(3nze)$
Stores only the non zero entries.
- Complicated data structure to assemble the matrix, especially in parallel.
- Can manage to apply different kinds of preconditioners.

$$A = \begin{bmatrix} 1. & 0. & 0. & 2. & 0. \\ 3. & 4. & 0. & 5. & 0. \\ 6. & 0. & 7. & 8. & 9. \\ 0. & 0. & 10. & 11. & 0. \\ 0. & 0. & 0. & 0. & 12. \end{bmatrix}$$



Conjugate Gradient Iterative Algorithm

- $$A\underline{x} = \underline{b}$$
1. Choose \underline{x}^0 , $\underline{r}^0 = \underline{b} - A\underline{x}^0$
 2. Solve $M\underline{z}^0 = \underline{r}^0$, $\underline{p}^0 = \underline{z}^0$
 3. $\gamma^0 = \text{Inner Product } (\underline{z}^0, \underline{r}^0)$
 - Do $k = 0, 1, 2, \dots k_{max}$
 4. $\underline{q}^k = \text{Matrix Multiply } (A, \underline{p}^k)$ ← Matrix-Vector Multiplication (1)
 5. $\tau^k = \text{Inner Product } (\underline{p}^k, \underline{q}^k)$ ← Inner Products (2)
 6. $\alpha^k = \gamma^k / \tau^k$
 7. $\underline{x}^{k+1} = \underline{x}^k + \alpha^k \underline{p}^k$, $\underline{r}^{k+1} = \underline{r}^k - \alpha^k \underline{q}^k$ ← Vector update (3)
 8. Solve $M\underline{z}^{k+1} = \underline{r}^{k+1}$ ← Preconditioner Solve (1)
 9. $\gamma^{k+1} = \text{Inner Product } (\underline{z}^{k+1}, \underline{r}^{k+1})$
 10. If $(\sqrt{\gamma^{k+1}} \leq \text{Tolerance})$ Exit
 11. $\beta^k = \gamma^{k+1} / \gamma^k$, $\underline{p}^{k+1} = \underline{z}^{k+1} + \beta^k \underline{p}^k$
 - End Do
- Note:** Independent of storage format for matrix



Reference: **Saad Yousef**, *Iterative Methods for Sparse Linear Systems*

Computational Mechanics Laboratory

CLEMSON
UNIVERSITY


Parallel Solution

To solve the FEA problem in parallel

- Form the system of equations in parallel.
- Store the matrix in distributed manner.
- Solve equations using parallel version of CG.


Steps :

- Mesh the domain.
- Decompose the domain into subdomains (METIS).
- Assign one processor to each subdomain.
- Each processor performs calculation for local elements (CSR requires assembly).
- Interior nodes are local to the processor.
- Interface nodes on subdomain boundaries requires communication (Update).



Domain Decomposition

The interior nodes are numbered first followed by each of the nodes lying on the subdomain interface. The system of equation then becomes:



Example mesh with 4 subdomains

$$\begin{bmatrix} A_1 & & & C_1 \\ & A_2 & & C_2 \\ & & \ddots & \vdots \\ & & & A_s & C_s \\ C_1^T & C_2^T & \cdots & C_s^T & A_I \end{bmatrix} \begin{Bmatrix} \underline{x}_1 \\ \underline{x}_2 \\ \vdots \\ \underline{x}_s \\ \underline{x}_I \end{Bmatrix} = \begin{Bmatrix} \underline{b}_1 \\ \underline{b}_2 \\ \vdots \\ \underline{b}_s \\ \underline{b}_I \end{Bmatrix}$$

$$A_I = \sum_{j=1}^s A_{I_j} \quad \text{and} \quad b_I = \sum_{j=1}^s b_{I_j}$$



Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

Block Matrix Structure

For each Process:

$$\begin{bmatrix} A_j & B_j \\ B_j^T & A_{I_j} \end{bmatrix} \begin{Bmatrix} \underline{x}_j \\ \underline{x}_{I_j} \end{Bmatrix} = \begin{Bmatrix} \underline{b}_j \\ \underline{b}_{I_j} \end{Bmatrix}$$

where,

A_j - Assembled local interior node matrix

A_{I_j} - Local interface nodes matrix

B_j - Coupling of the local interior nodes to local interface nodes


B_j^T - Coupling of the local interface nodes to local interior nodes

\underline{x}_j - Local interior solution subvector

\underline{x}_{I_j} - Local interface solution subvector

\underline{b}_j - Local interior force subvector

\underline{b}_{I_j} - Local interface force subvector



Subroutine Update (\underline{a})

```

    . . . Send . . . . .
    Do j = 1, np
        Buf(j) = a(values shared by j and this process)
        MPI Send (Buf, j)
    End do
    ----- Receive -----
    Do j = 1, np
        MPI Receive (Buf, j)
        a = a + Buf
    End do

```

End Subroutine Update



Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

Parallel Conjugate Gradient Algorithm

For each process j ;

1. $\underline{b}_{I_j} = \text{Update } (\underline{b}_{I_j})$

Requires pt. to pt. Communication
(Update).

2. $\underline{x}_j^0 = \underline{0}, \quad \underline{x}_{I_j}^0 = \underline{0}; \quad \underline{r}_j^0 = \underline{b}_j, \quad \underline{r}_{I_j}^0 = \underline{b}_{I_j}$

3. Solve $M(\underline{z}_j^0, \underline{z}_{I_j}^0) = (\underline{r}_j^0, \underline{r}_{I_j}^0); \quad \underline{p}_j^0 = \underline{z}_j^0, \quad \underline{p}_{I_j}^0 = \underline{z}_{I_j}^0$

4. $\gamma^0 = \text{Inner Product } (\underline{z}_j^0, \underline{z}_{I_j}^0; \underline{r}_j^0, \underline{r}_{I_j}^0)$

Do $k = 0, 1, 2, \dots$, Maximum Iterations

5. $(\underline{q}_j^k, \underline{q}_{I_j}^k) = \text{Matrix Multiply } (A, (\underline{p}_j^k, \underline{p}_{I_j}^k))$

Requires Update

6. $\tau^k = \text{Inner Product } (\underline{p}_j^0, \underline{p}_{I_j}^0; \underline{q}_j^0, \underline{q}_{I_j}^0)$

Requires one collective
communication

7. $\alpha^k = \gamma^k / \tau^k$

8. $\underline{x}_j^{k+1} = \underline{r}_j^k + \alpha^k \underline{p}_j^k; \quad \underline{x}_{I_j}^{k+1} = \underline{x}_{I_j}^k + \alpha^k \underline{p}_{I_j}^k$

9. $\underline{r}_j^{k+1} = \underline{r}_j^k - \alpha^k \underline{q}_j^k; \quad \underline{r}_{I_j}^{k+1} = \underline{r}_{I_j}^k - \alpha^k \underline{q}_{I_j}^k$

10. Solve $M(\underline{z}_j^{k+1}, \underline{z}_{I_j}^{k+1}) = (\underline{r}_j^{k+1}, \underline{r}_{I_j}^{k+1})$

11. $\gamma^{k+1} = \text{Inner Product } (\underline{z}_j^{k+1}, \underline{z}_{I_j}^{k+1}; \underline{r}_j^{k+1}, \underline{r}_{I_j}^{k+1})$

12. If $(\sqrt{\gamma^{k+1}} \leq \text{Tolerance})$ End

13. $\beta^k = \gamma^{k+1} / \gamma^k; \quad \underline{p}_j^{k+1} = \underline{z}_j^{k+1} + \beta^k \underline{p}_j^k, \quad \underline{p}_{I_j}^{k+1} = \underline{z}_{I_j}^{k+1} + \beta^k \underline{p}_{I_j}^k$

End Do



Reference: **Jimack P.K, Touheed N**, *Developing Parallel Finite Element Software Using MPI*
Computational Mechanics Laboratory



Matrix-by-Vector Multiply

$$\begin{Bmatrix} \underline{q}_j \\ \underline{q}_{I_j} \end{Bmatrix} = \begin{bmatrix} A_j & B_j \\ B_j^T & A_{I_j} \end{bmatrix} \begin{Bmatrix} \underline{p}_j \\ \underline{p}_{I_j} \end{Bmatrix} \quad j = 1, 2, \dots, np$$

For each process j :

1. $\underline{q}_j = A_j \underline{p}_j + B_j \underline{p}_{I_j}$
2. $\underline{q}_{I_j} = B_j^T \underline{p}_j + A_{I_j} \underline{p}_{I_j}$
3. Update (\underline{q}_{I_j})

Subroutine EBEmatmul $(Ke, p_j, p_{I_j}, q_j, q_{I_j})$

Do $e = 1, el$

Do $i = 1, nen$

If {Interior node}

Assemble the contributions to \underline{q}_j

Else If {Interface node}

Assemble the contributions to \underline{q}_{I_j}

End If

End do

End do

Update (\underline{q}_{I_j})

End Subroutine EBEmatmul

Subroutine CSRmatmul $(Acsr, p_j, p_{I_j}, q_j, q_{I_j})$

$\underline{q}_j = A_j p_j + B_j p_{I_j}$

$\underline{q}_{I_j} = B_j^T p_j + A_{I_j} p_{I_j}$

Send \underline{q}_{I_j} using MPI Send

Receive \underline{q}_j using MPI Recv

Add \underline{q}_{I_j}

End Subroutine CSRmatmul



Inner Product and Diagonal Preconditioner

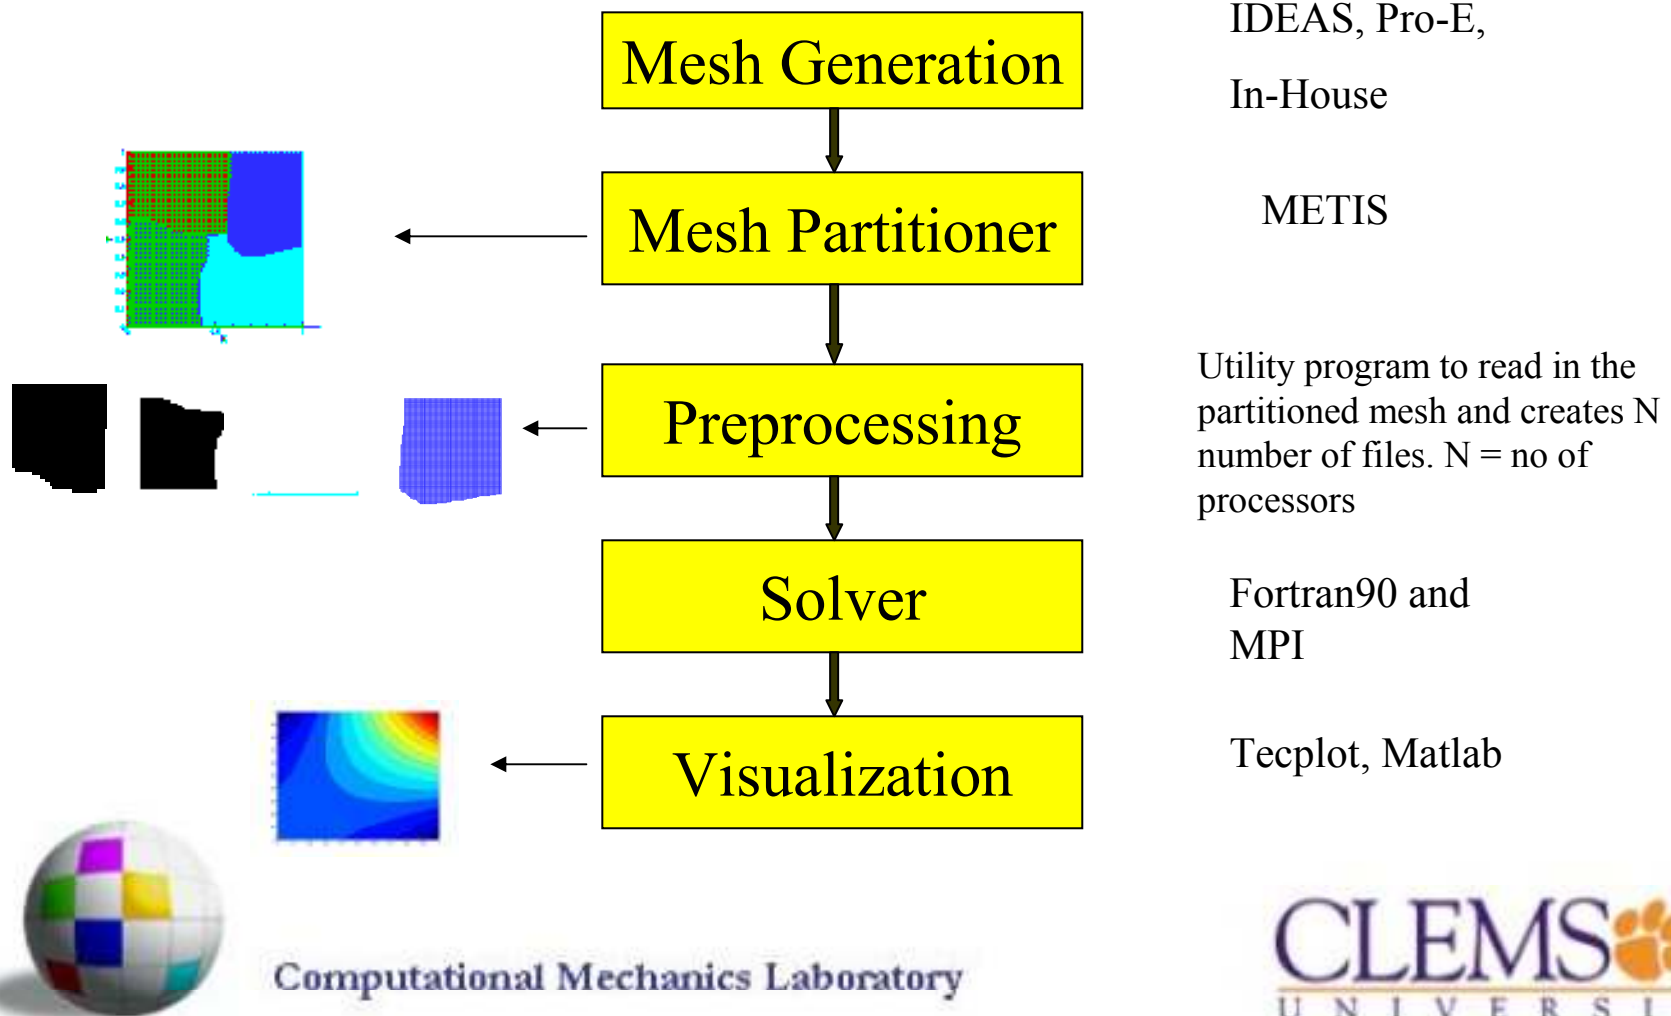
```
Function InnerProduct ( $\underline{a}_j, \underline{a}_{I_j}; \underline{b}_j, \underline{b}_{I_j}$ )
    sum = Dot_Product ( $\underline{a}_j, \underline{b}_j$ )
    n = size( $\underline{b}_{I_j}$ )
    sum = sum +  $\sum_{i=1}^n [\underline{a}_{I_j}(i) \times \underline{b}_{I_j}(i)] / \text{Shared}(i)$ 
    MPI_AllReduce( sum, InnerProduct, MPI_SUM)
End Function InnerProduct
```

```
For all process  $j$  ;
Subroutine SetupPreconditioner()
    Do  $e = 1, \dots, el$ 
        Assemble diagonal of  $A_j$  to  $\underline{d}_j$ 
        Assemble diagonal of  $A_{I_j}$  to  $\underline{d}_{I_j}$ 
        Update ( $\underline{d}_{I_j}$ )
        Invert the diagonal vectors  $\underline{d}_j = 1/\underline{d}_j$ ;     $\underline{d}_{I_j} = 1/\underline{d}_{I_j}$ 
    End Subroutine SetupPreconditioner
```

```
Subroutine Solve ( $\underline{z}_j, \underline{z}_{I_j}; \underline{r}_j, \underline{r}_{I_j}$ )
     $\underline{z}_j = \underline{d}_j * \underline{r}_j$ 
     $\underline{z}_{I_j} = \underline{d}_{I_j} * \underline{r}_{I_j}$ 
End Subroutine Solve
```



Implementation



Numerical Examples

Poisson's Equation with Dirichlet BC

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} = 0, \quad 0 < x < 1, \quad 0 < y < 1$$

Subject to Dirichlet boundary condition



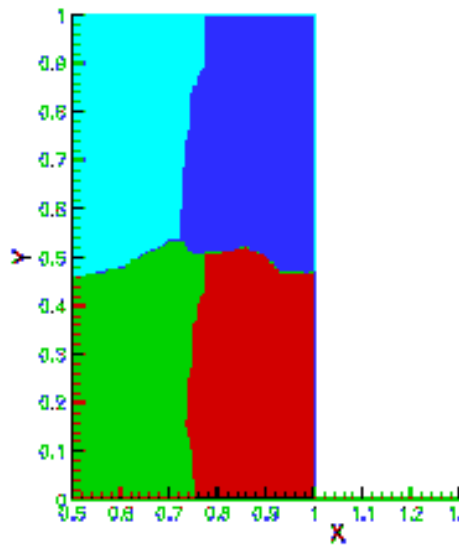
Exact solution is : $T(x, y) = \frac{\sin \pi x \cdot \sinh \pi y}{\sinh \pi}$

$$T(x = 0) = 0$$

$$T(x = 1) = 0$$

$$T(y = 0) = 0$$

$$T(y = 1) = \sin \pi x$$



Code Validation

Dimension of the mesh : 200 x 400

Number of DOF: 79,800

Tolerance : 1e-06



Parallel solution converges to exact solution.

Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

Analysis

2D structured mesh: 600 x 600 4-node Quads, 359,400 DOF

Implemented on

➤ Sun HPC

14 UltraSPARCII CPU's

336MHz

4MB cache

4GB shared memory


➤ SGI Onyx2 Infinite Reality system

8 MIPS R12000 CPU's

400 MHz

8MB cache

8GB shared memory



Convergence of CG and PCG.




Computational Mechanics Laboratory

CLEMSON
UNIVERSITY


Communication & Computation Time

SUN Machine

EBE



CSR




- Communication time increases with increase in number of processors.
- Computation time decreases with the increase in number of processors.




SGI Machine

EBE



CSR



- Communication time is virtually nil.




Computational Mechanics Laboratory

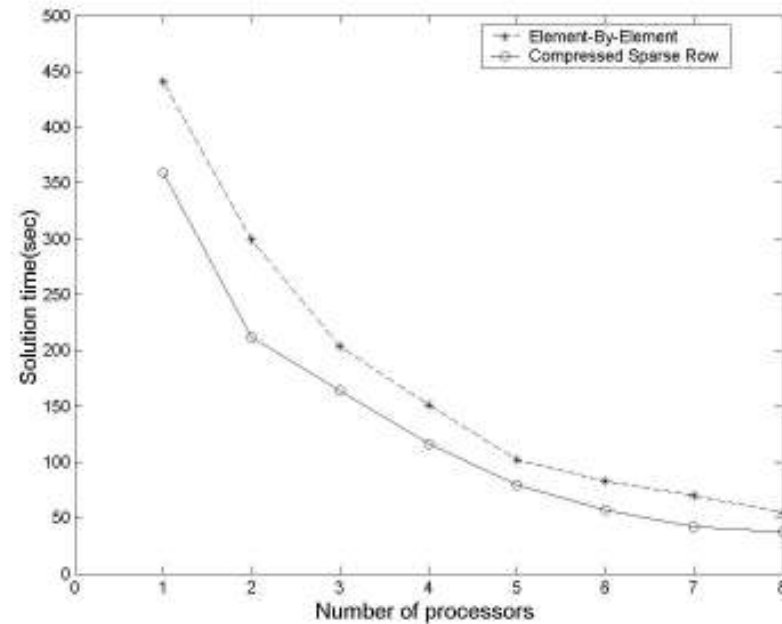
CLEMSON
UNIVERSITY

Solution Timings

Sun Machine



SGI Machine



CSR format better than EBE on both machines by almost 30%.

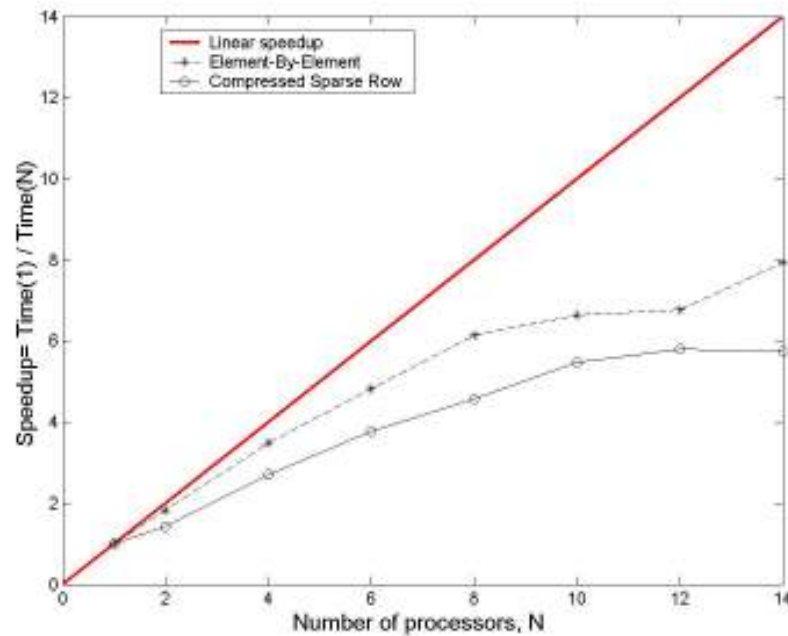


Computational Mechanics Laboratory


CLEMSON
UNIVERSITY

Speedup

Sun Machine



SGI machine



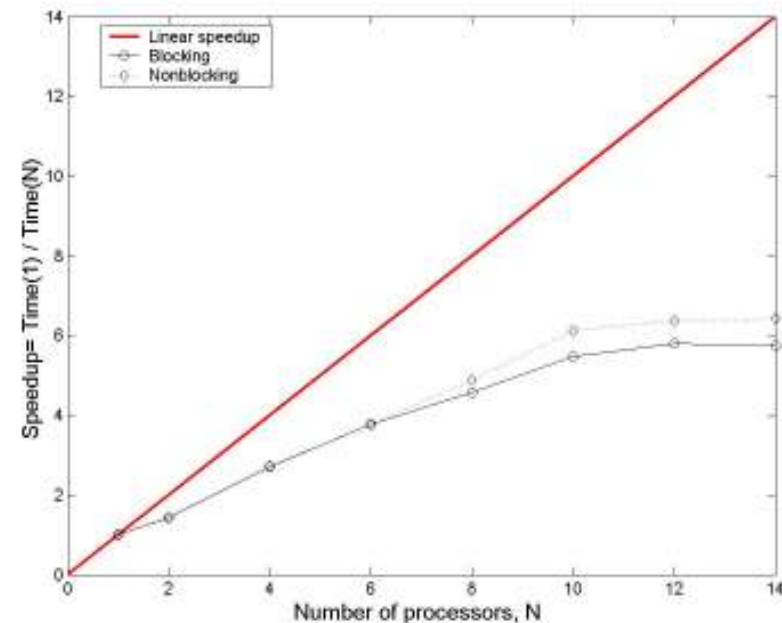

- Sun machine scales well up to 10 processors.
- SGI shows super-linear speedup for CSR.



Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

Comparison of Blocking and Nonblocking Communication



Nonblocking saves some communication time as the number of processors increase.

Subroutine CSRMatmul

$$\underline{q}_{I_j} = B_j^T \underline{p}_j + A_{I_j} \underline{p}_{I_j}$$

Send \underline{q}_{I_j} using MPIISend

$$\underline{q}_j = A_j \underline{p}_j + B_j \underline{p}_{I_j}$$

Receive \underline{q}_{I_j} using MPIRecv

Add \underline{q}_{I_j}


Subroutine CSRMatmul




Computational Mechanics Laboratory

CLEMSON
UNIVERSITY


Parallel code validation for unstructured Meshes – 2D



Unstructured Triangular Mesh
generated In I-DEAS




The Mesh is Partitioned
using METIS



Solved in Parallel using 4
Processor and Visualized
using Tecplot


I-DEAS
SOLUTION



Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

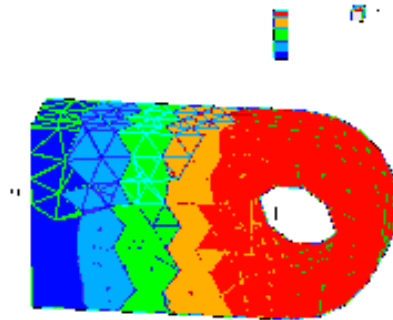
Parallel code validation for unstructured Meshes – 3D




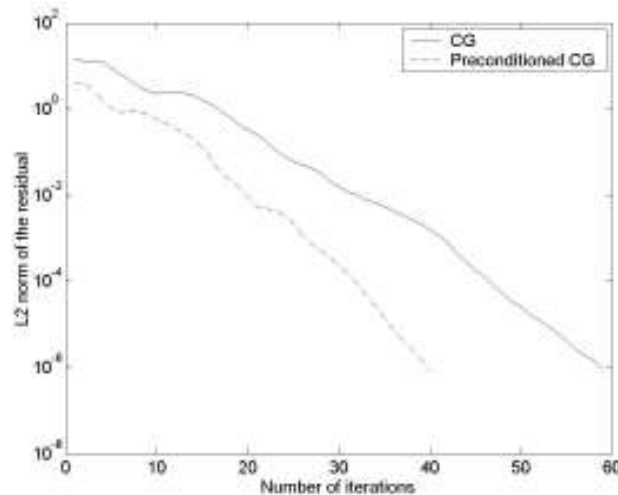
Unstructured Tetrahedra Mesh
generated In I-DEAS



The Mesh is Partitioned
using METIS



Solved in Parallel using 4
Processor and Visualized
using Tecplot



I-DEAS
SOLUTION



Computational Mechanics Laboratory

CLEMSON
UNIVERSITY

Conclusions

- CSR format takes 30% less CPU time than EBE.
- SGI with optimized vendor supplied MPI, shows super-linear speedup for the problem tested.
- Sun machine scales well up to 10 processors.
- By overlapping communication and computation total solution time can be reduced.
- Diagonal preconditioner moderately accelerates the convergence of the CG solver.



Future Work

- Perform scalability analysis on large 3D problems.
- Scalability analysis on Beowulf cluster.
- Implement Schur Complement method.
- Implement better preconditioner.
- Generalize to non-symmetric and non positive definite systems. Other iterative methods such as Bi-Conjugate Gradient Stabilized(Bi-CGS), QMR and GMRES.

